

THIS DOCUMENT COVERS

- ◆ [Positional Number Systems](#)
- ◆ [Computer Representation of Numbers](#)

Positional Number Systems

Overview

A positional number system represents any real number \Re as a polynomial in the base of the number system.

$$\pm(d_{\infty}\beta^{\infty} + \dots + d_1\beta^1 + d_0\beta^0 + d_{-1}\beta^{-1} + d_{-2}\beta^{-2} + \dots d_{-\alpha}\beta^{-\alpha}) = \pm\left(\sum_{k=-\infty}^{\infty} d_k\beta^k\right)$$

When writing polynomial representations of numbers we use a radix point to separate the whole and fractional parts. We can drop the powers of the base β as the exponent is implicit in the position of the digit. If a particular power has no value we still need to mark it with a co-efficient of zero. Our form becomes.

$$\pm(d_{\infty} \dots d_1 d_0 . d_{-1} d_{-2} \dots d_{-\infty})_{\beta}$$

The following are some examples

- $+34.15_{10} = (3 \times 10^1 + 4 \times 10^0 + 1 \times 10^{-1} + 5 \times 10^{-2})_{10}$
- $-11.01_2 = -(1 \times 2^1 + 1 \times 2^0 + 0 \times 2^{-1} + 1 \times 2^{-2})_2 = -3.25_{10}$

If we drop the fractional part of the representation we restrict ourselves to the set of integers \mathbb{Z} . In programming languages these are known as the ‘signed integer’ as they can be both positive and negative. If we restrict ourselves to only the positive whole numbers \mathbb{N} we have what programming languages call the ‘unsigned integers’

NOTE: \mathbb{N}

Mathematicians usually assume the natural numbers \mathbb{N} exclude zero. We use the standard computer science convention that the set \mathbb{N} includes zero.

Risk and Pricing Solutions

Normalized Scientific Form

Real numbers from the set \mathbb{R} form the basis of most scientific calculations. Any real number can be written in normalized scientific form.

$$mantissa \times \beta^e, 1.0 \leq mantissa < \beta, e \in \mathbb{Z}$$

The mantissa is a real number whose value is greater than or equal to 1.0 and less than β . The exponent is an integer. An example of a number in this form is

$$(3.1456224 \times 10^4)_{10}$$

Given an infinite number of digits in the fractional part of the mantissa any real number can be represented in this general form.

Restricting the representation size

In practice we do not have the luxury of an infinite number of digits and hence it is common to use a more restricted representation. In full generality, if we use a base β and have p digits of precision in the mantissa/significand we can **represent** a real number using the representation.

$$\mp (d_0.d_1d_2\dots d_{p-1}) \times \beta^e$$

This is the same as the following.

$$\mp (d_0\beta^0 + d_1\beta^{-1} + \dots + d_{p-1}\beta^{-(p-1)}) \times \beta^e, (1 \leq d_0 < \beta, 0 \leq d_{i=0..(p-1)} < \beta, e \in \mathbb{Z})$$

We also need to restrict the values of the exponent in β^e such that

$$(e \in \mathbb{Z}, e_{min} \leq e \leq e_{max})$$

Risk and Pricing Solutions

Properties of the finite representation


We now consider the important properties of a finite representation

$$\mp(d_0.d_1d_2\dots d_{p-1}) \times \beta^e, (1 \leq d_0 < \beta, 0 \leq d_{i:=1..(p-1)} < \beta, e \in \mathbb{Z}, e_{\min} \leq e \leq e_{\max})$$

DISTINCT VALUES

How many different numbers can a normalised finite form represent? One key point of the standard form is that the digit d_0 before the decimal point must be in the set $[1, \beta - 1]$ where the digits $d_2\dots d_3$ can be zero $[0, \beta - 1]$. In our representation we can calculate the total number of different representable values as the product of

$$2 \times (\beta - 1) \times \beta^{(p-1)} \times (e_{\max} - e_{\min} + 1)$$



LARGEST AND SMALLEST VALUES


The largest representable value is $(\beta - \beta^{-(p-1)})\beta^{e_{\max}}$. The smallest representable value is $-(\beta - \beta^{-(p-1)})\beta^{e_{\max}}$. The smallest non-negative value is $1.0 \times \beta^{e_{\max}}$.

NEAREST VALUES

The distance between two nearest values in our representation depends on the particular value of the exponent and the smallest non-zero positive value. $\beta^e \times \beta^{-(p-1)}$ Consider the following case where we use base 2, 4 digits for the mantissa and 2 digits for the exponent.

$$\mp(d_0.d_1d_2d_3) \times 2^e$$

If the exponent is 2^0 then our number line becomes



and the distance between the two nearest values is $2^{-3} \times 2^0 = 0.125_{10}$

Risk and Pricing Solutions

But if we increase the exponent to 2^1 our number line becomes



And the distance between the two nearest values becomes $2^{-3} \times 2^1 = 0.250_{10}$ and in general for a given exponent 2^e and a given number of binary digits p in the mantissa the distance between the nearest two points in our representation is $2^e \times 2^{-(p-1)}$

SUMMARY

PROPERTIES OF FINITE REAL NUMBER REPRESENTATIONS

◆ Possible different values	$2 \times (\beta - 1) \times \beta^{(p-1)} \times (e_{max} - e_{min} + 1)$
◆ Smallest non-zero positive value	$1.0 \times \beta^{e_{min}}$
◆ Largest representable value	$(\beta - \beta^{-(p-1)})\beta^{e_{max}}$
◆ Smallest representable value	$-(\beta - \beta^{-(p-1)})\beta^{e_{max}}$
◆ Difference between nearest 2 values	$\beta^e \times \beta^{-(p-1)}$

EXAMPLES

If we use base 2 with 4 digits for the mantissa and 2 digits for the exponent our finite normalized scientific representation becomes

$$\mp(d_0.d_1d_2d_3) \times 2^e$$

The leading integer digit of the mantissa must be non-zero in normalized notation and such a binary digit is in the set $[0,1]$ the only valid value it can take is 1. All the other digits in the mantissa and exponent can be either 0 or 1 giving us a total number of representable values as the product of

◆ 1	values of the integer part of the mantissa
◆ 2^3	values of the fractional part of the mantissa
◆ 2^2	values of the exponent
◆ 2	positive and negative values of the exponent
◆ 2	positive and negative values of the mantissa

Giving a total number of representable values of

Risk and Pricing Solutions

$$[1 \times 2^3 \times 2^2 \times 2 \times 2] = 128$$

We note an important point here. We used 4 bits for the mantissa and 2 bits for the exponent, one bit for the sign of the mantissa and one bit for the sign of the exponent coming to a total of 8 bits. However the total number of representable values is only $128 = 2^7$. This is because the leading integer digit of the mantissa has to be one. (remember in normalized scientific notation the integer digit must be greater than or equal to one and less than β . If β is 2 then only the integer digit 1 meets this criteria). We need one bit less in the representation. When we look at computer representation of floating point numbers later we will meet this again.

EXAMPLE 1 BASE $\beta = 10, p = 3, e_{max} = 99, e_{min} = -99$

- ◆ Smallest non-zero positive value 1.0×10^{-99}
- ◆ Largest representable value $(10 - 10^{-2})10^{99} = 9.99 \times 10^{99}$
- ◆ Smallest representable value $-(10 - 10^{-2})10^{99} = -9.99 \times 10^{99}$
- ◆ Difference between nearest 2 values $\beta^e \times \beta^{-(p-1)}$

EXAMPLE 2 BASE $\beta = 2, p = 2, e_{max} = 1, e_{min} = -1$

- ◆ Smallest non-zero positive value $(1.0 \times 2^{-1})_2 = (0.5)_{10}$
- ◆ Largest representable value $(1.1 \times 2^1)_2 = (1.5)_{10}$
- ◆ Smallest representable value $(-1.1 \times 2^1)_2 = (-1.5)_{10}$
- ◆ Difference between nearest 2 values

Risk and Pricing Solutions

Representation error

ABSOLUTE ERROR

Most real numbers can only be approximated by a finite representation. As such we need to be able to measure the error in approximation. If we are approximating some real number x with a floating point representation $\text{float}(x)$ the absolute error in the approximation is given by.

$$\text{Absolute error} = \text{float}(x) - x$$

RELATIVE ERROR

One problem with absolute error is that it does not take into account the scale of the number being approximated. Relative error includes the magnitude of the value we are approximating.

$$\text{Relative error} = \left| \frac{\text{float}(x) - x}{x} \right|$$

In the previous section we showed that the distance between the nearest representable values in a representation with p digits in the mantissa is $\beta^e \times \beta^{-(p-1)}$ for a particular value of the exponent e .

UNITS OF THE LAST PLACE (ULPS)

Often we are interested in the absolute error in terms of the precision of the mantissa, ignoring the exponent part. We often talk of the error in “units of the last place” which mean the error in units of the last place of the mantissa. So if our mantissa has 3 decimal digits in the fractional part then if our floating point representation of 5.413298 is given by 5.413 then the error in units of the last place would be .298. The unit of the last place itself has value $\beta^{-(p-1)}$

In our general form where we approximate a number x by a floating point number $\text{float}(x)$ with base β and p digits the error in units in the last place becomes.

$$\left| d_0.d_1d_2\dots d_{p-1} - \frac{x}{\beta^e} \right| \frac{1}{\beta^{-(p-1)}}$$

Risk and Pricing Solutions

Of course given a number in units of the last place $\beta^{-(p-1)}$ we simply multiply by the exponent term to get the absolute error

$$\text{absolute error} = \text{ulps} \times \beta^e = \beta^{-(p-1)} \beta^e = \beta \times \beta^{-p} \beta^e$$

If we have procedure that guarantees that the floating point number chosen to approximate our real number x is the closest floating point number then the error in terms of “units of the last place” can be at most $\frac{1}{2}$ times the value of the unit of the last place

$$\frac{1}{2} \times \beta^{-(p-1)} = \frac{1}{2} \times \beta^{-p+1} = \frac{\beta}{2} \times \beta^{-p}$$

$$0.5 \text{ulps} = \frac{\beta}{2} \beta^{-p}$$

FROM UNITS OF THE LAST PLACE TO RELATIVE ERROR

For any chosen value of e a number of the form $\mp(d_0.d_1d_2\dots d_{p-1}) \times \beta^e$ can vary in value from $\mp(1.00\dots 0) \times \beta^e$ all the way to $\mp((\beta - 1).(\beta - 1)(\beta - 1)\dots(\beta - 1)) \times \beta^e \approx \beta^e \times \beta$. As such for any chosen value of e , where the error in terms of ulps is fixed the relative error will vary from $\frac{\text{ulps} \times \beta^e}{\beta^e \beta}$ up to $\frac{\text{ulps} \times \beta^e}{\beta^e}$. If we have chosen the nearest floating point value to the real value then we can say that the error measured in ulps is 0.5 then our relative error will vary from $\frac{1}{2} \beta^{-p}$ up to $\frac{\beta}{2} \beta^{-p}$.

Put another way. For a fixed value of error in ulps the relative error can vary by a factor of β due to the fact that the mantissa can vary from 1.0 up to just under $\beta - B^{-p} \approx B$.

Risk and Pricing Solutions

We now consider a numerical example to cement these formulas. Consider the special case where we use a base of 10 and mantissa of 4 digits. Assuming we always choose the correct nearest floating point number then the error in ulps will be 0.5. In our case the last place has value 10^{-3} . Our chosen value of e is 3 so from our ulp error to absolute error we multiply 0.5×10^{-3} by 10^3 giving us an absolute error of 0.5 units. If we consider the value

1.000×10^3 our relative error becomes $\frac{0.5 \times 10^{-3} \times 10^3}{1.000 \times 10^3} = 0.5 \times 10^{-3} = \frac{1}{2} \beta^{-p}$ On the other hand

if we consider the value 9.999×10^3 and our relative error becomes $\frac{0.5 \times 10^{-3} \times 10^3}{9.999 \times 10^3} =$

$0.5 \times 10^{-2} = \frac{\beta}{2} \beta^{-p}$ Both of these confirm what we expect. A fixed absolute ulp for a given exponent e gives a relative error that varies by a factor of B depending on the value of the mantissa

Risk and Pricing Solutions

Division

Division is nothing but repeated subtraction. Integer division is defined using the following terms.

$$\text{dividend} = (\text{quotient} \times \text{divisor}) + \text{remainder}$$

And from this we have

$$\text{remainder} = \text{dividend} - (\text{quotient} \times \text{divisor}) = \text{dividend} \% \text{divisor}$$

And

$$\text{quotient} = (\text{dividend} - \text{remainder}) / \text{divisor}$$

Example of division

$$18 = (16 \times 1) + 2$$

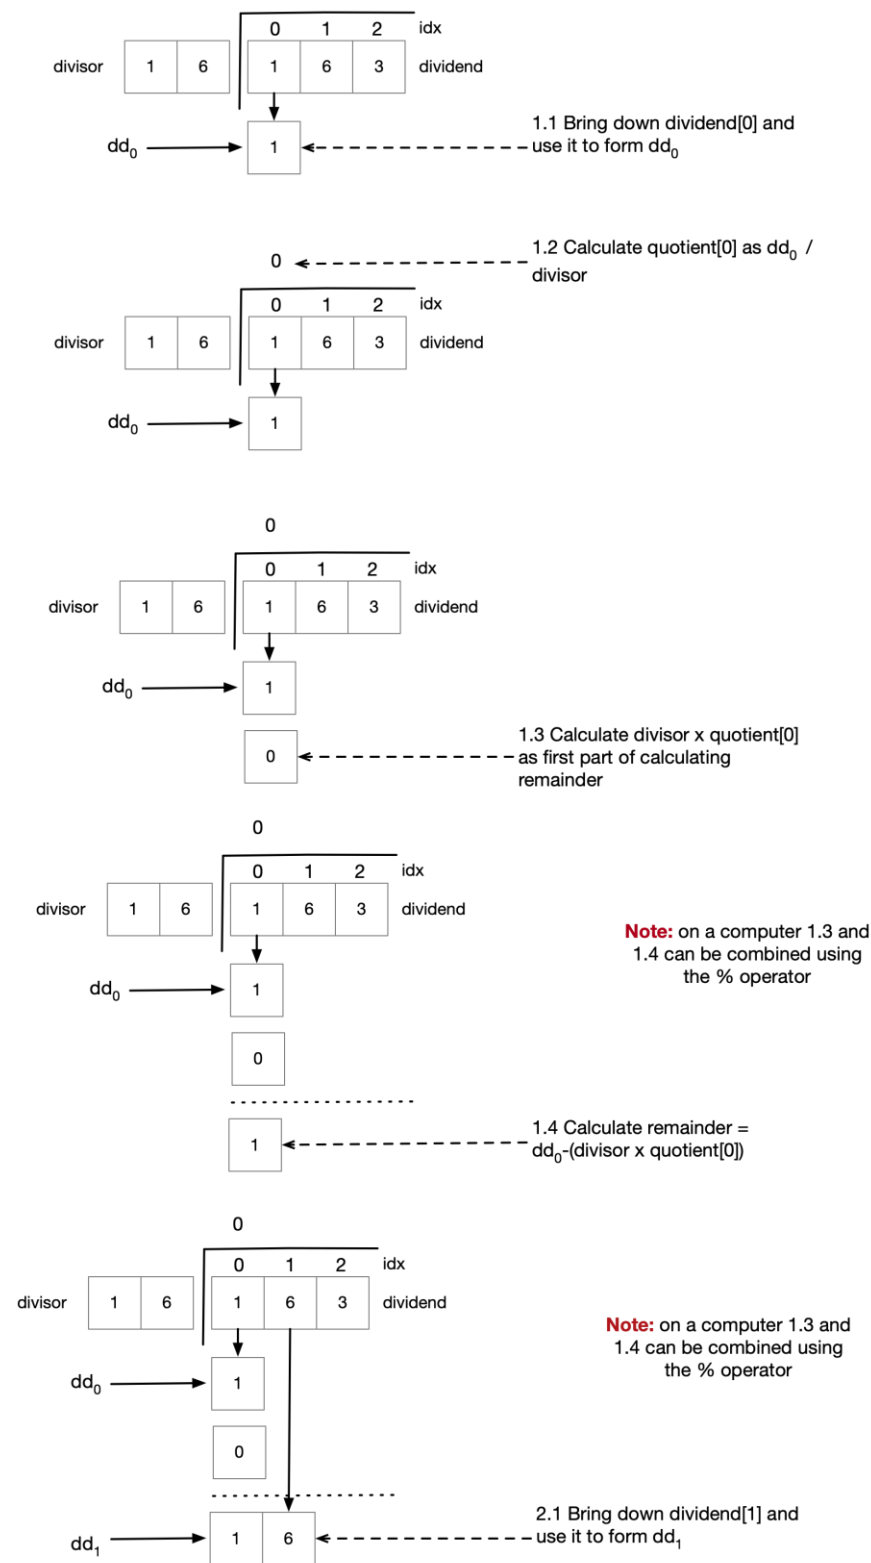
$$2 = 18 - (16 \times 1) = 18 \% 2$$

$$1 = 18 - \frac{(18 - 2)}{16} = 18 / 2$$

Risk and Pricing Solutions

Long Division

Consider $\frac{163}{16}$



Risk and Pricing Solutions

ALGORITHM - BASIC

```
private (string quotient, string remainder) DoLongDivision(string dividend, int
divisor, int b = 10)
{
    StringBuilder quotient = new StringBuilder();
    int remainder = 0;
    int dd = 0;

    for (int idx = 0; idx < dividend.Length; idx++)
    {
        // idx.1 copy in next digit into temporary dividend dd
        dd = (dd * b) + dividend[idx].ToIntDigit();

        // idx.2 calculate partial quotient and set into quotient[idx]
        int partialQuotient = dd / divisor;
        quotient.Append(partialQuotient.ToChar());

        // idx.3 calculate this temporary as part of calculating remainder
        int temp = partialQuotient * divisor;

        // idx.4 Calculate the remainder
        remainder = dd % divisor;

        // the remainder will form the basis of dd[idx+1]
        dd = remainder;
    }

    return (quotient.ToString(), remainder.ToChar().ToString());
}

public static class Extensions
{
    public static int ToIntDigit(this char c)
    {
        if (char.IsNumber(c)) return (int)char.GetNumericValue(c);

        return (int)(char.ToLower(c) - 'a' + 10);
    }

    public static char ToChar(this int i)
    {
        if (i >= 0 && i < 10)
            return (char)(i + '0');

        return (char)(i + 'a' - 10);
    }
}
```

Risk and Pricing Solutions

ALGORITHM – FLOATING POINT RESULTS

```
public string IntegerDivisionWithFloatingPointResult(string dividend, int divisor,
    int b = 10, int maxDigits = 8)
{
    StringBuilder quotient = new StringBuilder();
    int remainder = 1;
    int dd = 0;

    for (int idx = 0; (idx < dividend.Length || remainder > 0)
        && idx < maxDigits; idx++)
    {
        // Add in a decimal point
        if (idx == dividend.Length)
            quotient.Append(".");

        // idx.1 copy in next digit into temporary dividend dd
        if (idx < dividend.Length)
            dd = (dd * b) + dividend[idx].ToIntDigit();
        else
            // The integer dividend has no more digits so we just increase
            // by a factor of b as we move to the right side of the point
            // point
            dd = (dd * b);

        // idx.2 calculate partial quotient and set into quotient[idx]
        int partialQuotient = dd / divisor;
        quotient.Append(partialQuotient.ToChar());

        // idx.3 calculate this temporary as part of calculating remainder
        int temp = partialQuotient * divisor;

        // idx.4 Calculate the remainder
        remainder = dd % divisor;

        // the remainder will form the basis of dd[idx+1]
        dd = remainder;
    }

    return quotient.ToString();
}
```

Risk and Pricing Solutions

Converting between bases

Give a number N in base λ we want to convert it to a new base β . That is to say given

$$N = \pm(a_n\lambda^\infty + \dots + a_2\lambda^1 + a_1\lambda^0 + b_1\lambda^{-1} + b_2\lambda^{-2} + \dots b_\alpha\lambda^{-\alpha})_\lambda$$

We want to find the coefficients c_i and d_i such that

$$N = \pm(c_n\beta^\infty + \dots + c_2\beta^1 + c_1\beta^0 + d_1\beta^{-1} + d_2\beta^{-2} + \dots d_\alpha\beta^{-\alpha})_\beta$$

When doing the conversion we consider the integral and fractional part separately.

Integral part

Looking first at the integral part we have a number N

$$N = (a_n a_{n-1} \dots a_2 a_1)_\lambda$$

We want to convert it to base β such that

$$N = (c_m c_{m-1} \dots c_1 c_0)_\beta$$

We can rewrite this as

$$N = c_1 + \beta(c_2 + \beta(c_3 + \dots + \beta(c_m)) \dots)_\beta$$

If we divide it by β then the remainder is clearly c_1 and the quotient is

$$c_2 + \beta(c_3 + \beta(c_4 + \dots + \beta(c_m)) \dots)_\beta$$

If we repeat this until the quotient is zero we can read off the value of c_1 to c_n giving us the required number in the new base $(c_m c_{m-1} \dots c_1 c_0)_\beta$

Let us consider the scenario where we want to convert the decimal number 2748 to hexadecimal. We first divide our decimal number by the new base 16

Risk and Pricing Solutions

$$\begin{array}{r} 171 \\ 16 \overline{) 2748} \\ \underline{1600} \\ 1148 \\ \underline{112} \\ 28 \\ \underline{16} \\ 12 \end{array}$$

So after this first division we know that

$$1) \quad 2748 = [(171 \times 16)] + 12$$

We can't represent 171 as we only have sixteen symbols so we to divide 171 by 16

$$\begin{array}{r} 10 \\ 16 \overline{) 171} \\ \underline{160} \\ 11 \end{array}$$

So now we know that

$$2) \quad 171 = [(10 \times 16)] + 11$$

Inserting ii) into i) we get

$$3) \quad 2748 = [(\{[10 \times 16] + 11\} \times 16)] + 12 = (16^2 \times 10) + (16^1 \times 11) + (16^0 \times 12)$$

Which we know is a positional number ABC_{16}

Similarly we can do the same for base 2

Fractional part

Consider the situation where we have a fraction part $0 < x < 1$ in some base λ and we want to find the digits d_k in the representation

Risk and Pricing Solutions

$$x = \sum_{k=1}^{\infty} d_k \beta^{-k} = (0.d_1 d_2 d_3 \dots)_\beta$$

We first note that

$$\beta x = (d_1.d_2 d_3 \dots)_\beta$$

So if we take our fractional part and multiply it by β then the resulting integral component is the d_1 we can similarly repeat the process to find the digits $d_2 \dots d_m$

EXAMPLE 1 CONVERT 0.526_{10} TO BASE 8

i) $8 \times 0.526_{10} = 4.208 \therefore 0.526_{10} = \frac{1}{8}4 + \frac{1}{8}(0.208)$

$$8 \times 0.208_{10} = 1.664 \therefore 0.208_{10} = \frac{1}{8}1 + \frac{1}{8}(0.664)$$

$$8 \times 0.664_{10} = 5.312 \therefore 0.664_{10} = \frac{1}{8}5 + \frac{1}{8}(0.312)$$

$$8 \times 0.312_{10} = 2.496 \therefore 0.312_{10} = \frac{1}{8}2 + \frac{1}{8}(0.496)$$

$$0.526_{10} \approx 0.4152_8$$

Risk and Pricing Solutions

Computer Representation of numbers

Unsigned integers – the whole numbers

INTRODUCTION


Consider our positional number system

$$\pm(d_{\infty}\beta^{\infty}+\dots d_1\beta^1 + d_0\beta^0 + b_{-1}\beta^{-1} + b_{-1}\beta^{-2}+\dots b_{-\alpha}\beta^{-\alpha}) = \pm\left(\sum_{k=-\infty}^{\infty} d_k\beta^k\right)$$

If we only need to represent whole numbers, that is to say unsigned integers we can use an n-bit binary representation. We don't need any bits to represent fractions.

$$(d_{n-1}\dots d_1d_0)_2 = (d_{n-1}2^{n-1}+\dots d_12^1 + d_02^0) = \pm\left(\sum_{k=0}^{n-1} d_k2^k\right)$$


Such a representation can distinguish between 2^n different values. It is often useful to visualize the representation as a circle mod 2^n . In the special case where $n = 3$ we get.



If we use our 2^n different values to represent positive integers in the range $[0, 2^n - 1]$ we get the following.

Risk and Pricing Solutions

UNSIGNED BINARY INTEGER ADDITION



In this representation addition is simply moving clockwise around the circle. It can easily be achieved naturally using binary addition in hardware. We can easily simulate unsigned addition in code via the bitwise shift and logic operators.

$$\begin{array}{r} 010 \text{ (2)} \\ 011 \text{ (3)} \\ \hline 101 \text{ (5)} \end{array}$$

Risk and Pricing Solutions

C++ CODE TO PERFORM BINARY ADDITION ON UNSIGNED INTEGERS

```
short binaryAdd( short x, short y )
{
    short result = 0.0;
    short carry = 0.0;

    int numberOfBits = sizeof(x) * 8.0;

    for ( int bitNumber = 0; bitNumber < numberOfBits; bitNumber++ )
    {
        // We deal with one bit at a time. By right shifting
        // x and y bitNumber times we can set the bit we
        // want into the least significant bit
        // of the two's complement representation.
        short shiftedX = x >> bitNumber;
        short shiftedY = y >> bitNumber;

        // Now we make use of the fact that the number 1 in
        // two's complement has (numberOfBits - 1) zeros
        // followed by a solitary 1 in the least significant
        // digit. We can then take our shifted values and bitwise
        // and them with 1 to make sure the only digit in the
        // shifted number is the one we want to deal with
        short xDigit = shiftedX & 1;
        short yDigit = shiftedY & 1;

        // We have three values that feed into the current digit
        // {the x digit, the y digit, the carry digit}. If
        // one or all three of these are zero then the
        // digit value will be zero, otherwise it will be one
        short digitValue = ( xDigit ^ yDigit ) ^ carry;

        // We now shift the digit back into its correct
        // location and add it to the result we are building up
        result |= ( digitValue << bitNumber );

        // Finally calculate the carry for the next round
        carry = ( xDigit & yDigit ) | ( xDigit & carry ) | ( yDigit &
carry );
    }

    return result ;
}
```


Risk and Pricing Solutions

UNSIGNED BINARY INTEGER SUBTRACTIONS

We have quite easily implemented code to perform addition of unsigned integers using bit operators. We now consider subtraction of unsigned integers. Unsigned integer subtraction makes use of some interesting properties of our binary number system.

1) $a_{\text{mod } 2^n} + 2^n = a_{\text{mod } 2^n}$ in an n bit unsigned representation

Essentially we are just moving one complete revolution back to the same number.



The second result we need is that in modular arithmetic (see textbox for proof)

2) $a_{\text{mod } 2^n} - b_{\text{mod } 2^n} = a_{\text{mod } 2^n} + [(2^n - 1) - b_{\text{mod } 2^n}] + 1$

- i) $a_{\text{mod } 2^n} - b_{\text{mod } 2^n} = [a_{\text{mod } 2^n} - b_{\text{mod } 2^n}] + 2^n$ From property 1)
- ii) $[a_{\text{mod } 2^n} - b_{\text{mod } 2^n}] + 2^n = a_{\text{mod } 2^n} + [2^n - b_{\text{mod } 2^n}]$ Rearranging i)
- iii) $a_{\text{mod } 2^n} - b_{\text{mod } 2^n} = a_{\text{mod } 2^n} + [2^n - b_{\text{mod } 2^n}]$ Subst rhs ii) into rhs of i)
- iv) $2^n = [2^n - 1] + 1$ Basic arithmetic
- v) $a_{\text{mod } 2^n} - b_{\text{mod } 2^n} = a_{\text{mod } 2^n} + [(2^n - 1) + 1 - b_{\text{mod } 2^n}]$ Sub iv) into rhs of iii)
- vi) $a_{\text{mod } 2^n} - b_{\text{mod } 2^n} = a_{\text{mod } 2^n} + [(2^n - 1) - b_{\text{mod } 2^n}] + 1$ Take 1 outside the brackets

The final clever step to performing subtraction of unsigned integers is to note that in an n bit unsigned representation where the bitwise complement operator is \sim .

Risk and Pricing Solutions

$$3) \sim b = (2^n - 1) - b$$

Proof

w	0	1	1
~w	1	0	0
~w + w	1	1	1

$$b + \sim b = \therefore (2^n - 1) - b = \sim b$$

We can then substitute this into 2)

$$a_{\text{mod } 2^n} - b_{\text{mod } 2^n} = a_{\text{mod } 2^n} + \sim b + 1$$

This is very powerful. It means that if we couple the code we already wrote to do unsigned addition with the bitwise complement operator we can then do unsigned subtraction as well.

Figure 1 Addition of negative unsigned integers

```
short binaryNegate( short x )
{
    short neg = binaryAdd(~x, 1);
    return neg;
}

short binarySubtract( short x, short y )
{
    short minusY = binaryNegate(y);

    return binaryAdd( x, minusY);
}
```

Risk and Pricing Solutions

Signed Integers

TWOS COMPLEMENT

Twos complement is a way of encoding negative numbers into ordinary binary such that addition still works. In the previous section we discussed how a polynomial can be represented as

$$w = a_n R^{n-1} + \dots + a_2 R^1 + a_1 R^0$$

Specifically letting R be 2 we can represent an N bit binary number as

$$w = a_n 2^{n-1} + \dots + a_2 2^1 + a_1 2^0$$

Using summation notation we have


$$w = \sum_{i=1}^n a_n 2^{i-1}$$

So a three bit number 111_2 in this representation would equal 7

In a twos complement signed representation we change the most significant digits weighting to $-1 \times a_n 2^{n-1}$ giving us

$$w = -1 \times a_n 2^{n-1} + \sum_{i=1}^{n-1} a_n 2^{i-1}$$

So in a three bit 2's complement notation the number 111_2 would equal $-4 + 2 + 1 = -1$ and indeed for any value n where all the coefficients are set to 1 $1 \dots 1_2$ the value is equal to -1.



Risk and Pricing Solutions

0	1	1	3	$(-1 \times 0 \times 2^2) + (1 \times 2^1) + (1 \times 2^0) = 3$
0	1	0	2	
0	0	1	1	
0	0	0	0	
1	1	1	-1	
1	1	0	-2	
1	0	1	-3	$(-1 \times 1 \times 2^2) + (0 \times 2^1) + (1 \times 2^0) = -3$
1	0	0	-4	$(-1 \times 1 \times 2^2) + (0 \times 2^1) + (0 \times 2^0) = -4$

Remember in the previous section on unsigned integers we saw that the maximum value that can be represented in an n-bit unsigned integer is $(2^n - 1)$. We also proved that subtracting a value from the binary representation with the value one in every bit is the same as flipping the bits

$$(2^n - 1) - b = \sim b$$

In our twos complement notation the binary value with a one in every bit is no longer $(2^n - 1)$ but instead is -1. Our equation then becomes

$$-1 - b = \sim b$$

And so

$$-b = \sim b + 1$$

So given a positive twos complement number we need to flip its bits and add one.

WHY TWOS COMPLEMENT IS POWERFUL

The most powerful aspect of the two complement notation is that we can add positive and negative numbers. If we have n bits we can represent 2^n values and hence due to overflow if we move 2^n points around our modular system we get back to the same number.

The algorithm to multiply a twos complement number by -1 is to flip all its bits using the logical negation operator \sim and then add one. This is beautiful because we can use the same

Risk and Pricing Solutions

bitwise addition to perform addition and subtraction. To do subtraction just form the ones complement and then do normal addition

FIGURE 2 BINARY NEGATION

```
short binaryNegate( short x )
{
    short neg = binaryAdd(~x, 1);
    return neg;
}
```

SUMMARY

- ◆ The most significant bit represents the sign
- ◆ Negating a value requires switching all its bits and then adding one
- ◆ 1 is represented by 001 and -1 is represented by 111
- ◆ N-bit implementation can represent numbers from -2^{n-1} to $2^{n-1} - 1$

Risk and Pricing Solutions

Floating Point representations

“The exact meaning of single-, double-, and extended-precision is implementation-defined. Choosing the right precision for a problem where the choice matters requires significant understanding of floating point computation. If you don’t have that understanding, get advice, take the time to learn, or use double and hope for the best”

Bjarne Stroustrup – The C++ Programming Language

INTRODUCTION

In the above section we saw that we can represent any non-zero real number using the notation

$$\mp(d_0.d_1d_2\dots d_{p-1}) \times \beta^e, (1 \leq d_0 < \beta, 0 \leq d_{i=1..(p-1)} < \beta, e \in \mathbb{Z}, e_{min} \leq e \leq e_{max})$$

Internally real numbers are stored in binary representation, i.e. our base is 2.

$$\mp(d_0.d_1d_2\dots d_{p-1}) \times 2^e, (d_0 = 1, d_{i=1..(p-1)} \in \{0,1\})$$

The key aspects of this representation are

1. $d_0.d_1d_2\dots d_{p-1}$ is known as the significand or the mantissa
2. The number of digits in the mantissa/significand p
3. The max and minimum exponents e_{min} and e_{max}
4. The base β

All floating point numbers are **rational numbers** which means they have a terminating expansion in the relevant base. As such most real numbers cannot be expressed exactly. Any number with an infinite expansion cannot be represented.

Also a number which has a finite expansion in one base can have non-finite expansion in another base. If the base is 2, as in binary floating point only **rational** numbers whose denominators are powers of 2 can be represented.

$$\frac{1}{2}, \frac{1}{4}, \frac{3}{4}, \frac{1}{8}, \frac{3}{8}, \frac{5}{8}, \frac{7}{8}$$

Converting a base 10 fraction such as 0.1 to binary floating point will result in an infinite expansion which can only be approximated with a finite number of digits. The following section discusses how to measure the error in any approximation

Risk and Pricing Solutions

NEAREST VALUES

Consider a floating point representation with base $\beta = 10, p = 3, e \in \{1, 0, -1\}$ The difference between the nearest two numbers depends on the exponent e .

$1.01 \times 10^1 = 10.1$	0.1
$1.01 \times 10^1 = 10.2$	0.1
$1.01 \times 10^1 = 10.3$	0.1
$1.01 \times 10^0 = 1.01$	0.01
$1.01 \times 10^0 = 1.02$	0.01
$1.01 \times 10^0 = 1.03$	0.01
$1.01 \times 10^{-1} = 0.101$	0.001
$1.01 \times 10^{-1} = 0.102$	0.001
$1.01 \times 10^{-1} = 0.103$	0.001

Generalising, the spacing between the two nearest values for a particular e is given by

$$\beta^{-(p-1)}\beta^e = \beta^{e-(p-1)}$$

If we need to round a real number to the nearest machine representable number the maximum absolute error is half the space or $\frac{\beta^{e-(p-1)}}{2}$

UNITS IN THE LAST PLACE (ULP)

Consider the case where we wish to represent 42.314 using our finite representation. Shifting the number into normalised form we obtain $42.314 = 4.2314 \times 10^1$ The nearest representable value is $42.3 = 4.23 \times 10^1$ The absolute difference between the two numbers is

z	$4.2314 \times 10^1 = 42.314$
f	$4.23 \times 10^1 = 42.3$
$z - f$	$0.0014 \times 10^1 = 0.014$

Since one unit in the last place is 0.1 our absolute error of 0.014 is equal to 0.14 units in the last place. We can calculate the value of 1 ulp by taking $\beta^{-(p-1)}\beta^e = 10^{-(3-1)}10^1 = \mathbf{0.1}$

$$1ulp = \beta^{-(p-1)}\beta^e$$

If we approximate a real number z using the machine number $d_1.d_2d_3 \times 10^1$ then the error in units of the last place is given by $|d_1.d_2d_3 \times 10^1 - z|\beta^{p-1}$ In full generality, given a base of β , a precision of p and an exponent of e we can calculate the error in approximating a real number z with the nearest machine number $d_0.d_1\dots d_{p-1} \times \beta^e$, measure in ulps as

$$ulps = \left| d_1.d_2d_{p-1} - \frac{z}{\beta^e} \right| \beta^{p-1}$$

Risk and Pricing Solutions

When approximating a real number with the nearest floating point value the error can be as much as $\left(\frac{\beta}{2}\beta^{-p}\right)\beta^e$

RELATIVE ERROR

Using the example from the previous section the relative error is given as

$$\frac{3.1459 - 3.14}{3.1459} = 0.000506$$

Now to see the relationship between absolute and relative error consider approximating the following two real numbers with the nearest floating point representatives; 9.995 and 0.005 (As we are looking at relative error the magnitude given by β^e can be ignored) The relative error of the two numbers are

$$\frac{9.995 - 9.99}{9.995} = \frac{0.005}{9.995} = 0.0005$$

$$\frac{1.005 - 1.00}{1.005} = \frac{0.005}{1.005} = 0.005$$

So although all numbers with a given β^e have the same maximum absolute error, there relative error varies by a factor of β^e This is known as the wobble.

SINGLE PRECISION IEEE

If we consider single precision number $x = \mp q \times 10^m$ the valid values the 32 bits are allocated as

$$\mp(d_0.d_1d_2\dots d_{23}) \times 2^{e_8e_7\dots e_1}$$

- ◆ 1 bit represents the sign of the number \mp

Risk and Pricing Solutions

- ◆ 23 bits for the fractional part of the mantissa
- ◆ 8 bit signed number for the exponent

The storage however is a little peculiar. We might expect that using 8 bits for the exponent would allow use to have 256 different values. However four values are reserved for special values such as plus and minus zero and plus and minus infinity.

The representation is $(-1)^s \times 2^{c-127} \times (1.f)_2$ where $-126 \leq c \leq 127$ (0 and 255 are used for special values) and $1 \leq (1.f)_2 \leq (1.111111111111111111111111)_2 = 2 - 2^{-23}$. The largest possible value representable is hence $(2 - 2^{-23})2^{127} \approx 2^{128} \approx 3.4 \times 10^{38}$. The smallest positive number becomes $(1)2^{-126} \approx 1.2 \times 10^{-38}$.

The binary machine number $\varepsilon = 2^{-23}$ is the machine epsilon and is hence the smallest positive value such that $1 + \varepsilon \neq 1$. Because $2^{-23} \approx 1.2 \times 10^{-7}$ we can infer that single precision floating point has accuracy to six significant decimal figures.

So the mantissa can represent from 1 to $2 - 2^{-23}$ in increments of 2^{-23} which in decimal is approximately from 1 to $2 - 1.2 \times 10^{-7}$ in increments of 1.2×10^{-7} so since any single precision mantissa representation can be up to 1.2×10^{-7} from the real number. As such the precision of the mantissa is 6 significant figures.

Questions

What is the precision of a single precision point floating point number and why?

Six significant figures

The binary machine number $\varepsilon = 2^{-23}$ is the machine epsilon and is hence the smallest positive value such that $1 + \varepsilon \neq 1$. Because $2^{-23} \approx 1.2 \times 10^{-7}$ which if we write it out we see

0.00000012 If we see this value what it really means is that the value is

$$0.00000018 > x > 0.00000006$$

So only the sixth significant figure is accurate.

What is the range of a single precision floating point and why?

From $\approx 2^{128}$ to $\approx -(2^{128})$ which is approximately from 3.4×10^{38} to $-(3.4 \times 10^{38})$

The reason being that the largest absolute value representable in single precision is given by $(2 - 2^{-23})2^{127} \approx 2^{128} \approx 3.4 \times 10^{38}$ as the mantissa has 23 bits and the exponent has 8 bits.

What is the precision of a double precision point floating point number and why?

The binary machine number $\varepsilon = 2^{-53}$ is the machine epsilon and is hence the smallest positive value such that $1 + \varepsilon \neq 1$. Because $2^{-53} \approx 1.1 \times 10^{-16}$ so only to the 15 significant figure is correct.

Given an operator that does addition how can one add subtract one unsigned integer from the other without using the subtraction operator?

$$x - b = x + \sim b + 1$$

Why does this work?

If we use n bits to represent unsigned integers addition is modulo 2^n giving a maximum value of $2^n - 1$

$$x + 2^n = x$$

$$x - b = x + 2^n - b$$

$$x - b = x + (2^n - 1) + 1 - b$$

Risk and Pricing Solutions

But $2^n - 1$ consists of n ones 1111...11 Subtracting from n ones is equivalent to flipping the bits because $1 - 0 = 1$ and $1 - 1 = 0$. So we can replace $(2^n - 1) - b = \sim b$ Substituting back in we get

$$x - b = x + \sim b + 1$$

What is the result of the following C code and why?

```
int fraction = 5 / 9 ;
```

The result is that the value 0 is assigned to fraction because in C integer division truncates.